


# Dask integration

# About Dask

- Transparently handles the brokerage of large datasets
- Transparently handles the distribution to workers
- Two relevant approaches:
  - “**processes**” (using multiprocessing.Pool —> iqqomp)
    - Over ~80% of the use cases
  - “**distributed**” (needs to define cluster object and pass it to Dask)
    - Less than ~20% of the use cases
    - e.g. dask-yarn handles with clusters but we don't need this extra layer

# About Dask




# Attempt with “processes”

- Naively we could use iqqomp instead of multiprocessing
- Problem is that desk has it builtin internally

## Native multiprocessing+desk (runs out of the box...)

```
1 import desk
2 from multiprocessing.pool import Pool
3 import os
4 from desk import compute, delayed
5
6 desk.config.set(pool=Pool(5))
7
8 def do_something(x): return x * x
9
10 data = range(100000)
11 delayed_values = [delayed(do_something)(x) for x in data]
12 results = compute(*delayed_values, scheduler='processes')
```



## iqqomp+desk (fails - see next slide)

```
1 import desk
2 # from multiprocessing.pool import Pool
3 import os
4 from iqqomp.pool import Pool
5 from desk import compute, delayed
6
7 os.environ['IQOQO_LOGIN_USER'] = 'efrat.tal@iqogo.co'
8 os.environ['IQOQO_LOGIN_PASSWORD'] = '12345678'
9
10 desk.config.set(pool=Pool(5))
11
12 def do_something(x): return x * x
13
14 data = range(100000)
15 delayed_values = [delayed(do_something)(x) for x in data]
16 results = compute(*delayed_values, scheduler='processes')
```

# Dask “processes”

```
→ Documents python3 dask_test.py
Traceback (most recent call last):
  File "dask_test.py", line 18, in <module>
    results = compute(*delayed_values, scheduler='processes')
  File "/usr/local/lib/python3.7/site-packages/dask/base.py", line 446, in compute
    results = schedule(dsk, keys, **kwargs)
  File "/usr/local/lib/python3.7/site-packages/dask/multiprocessing.py", line 199, in get
    len(pool._pool),
AttributeError: 'Pool' object has no attribute '_pool'
```

- This image just shows the beginning of the process, i.e. even if I add the `_pool` object to the class, things fail downstream
- Will need to hack Dask or \*fully\* implement `iqoqomp` Pool class as the Pool class of `multiprocessing.py` (at least)
- Not shown here but also the multithreaded option is impossible

# iqoqomp.Pool

```
76     def apply(self, func, args=(), kwds={}):
77         raise NotImplementedError
78
79     def apply_async(self, func, args=(), kwds={}, callback=None,
80                    error_callback=None):
81         raise NotImplementedError
82
83     def map_async(self, func, iterable, chunksize=None, callback=None,
84                  error_callback=None):
85         raise NotImplementedError
86
87     def starmap_async(self, func, iterable, chunksize=None, callback=None,
88                      error_callback=None):
89         raise NotImplementedError
90
91     def imap(self, func, iterable, chunksize=1):
92         raise NotImplementedError
93
94     def imap_unordered(self, func, iterable, chunksize=1):
95         raise NotImplementedError
```

- Many features are not implemented (on purpose)
- Using “processes” approach, Dask.compute() is expecting a complete Pool object and it relies on the internal copy of multiprocessing which has much more functionalities
  - This approach is really single-machine-oriented and many things are hardcoded there for that reason
  - It could be solved with some adaptations on their side, **but** this is why they have the “distributed” architecture, so I doubt they will accept any change we may come with

# “distributed”

- It is not the most frequent use case because it is “more complex” for random users
- Need direct access to the cluster
  - Dask provides some interfaces with Amazon / Google /... clouds
  - Example with Yarn below (irrelevant for us but just to make a point)
    - It provides the entire machinery that we provide, including the definition of the cluster
- If we go that way, we need to write a wrapper that returns the iqoqo “cluster” as a one-liner

**Estimate the implementation to require ~4 weeks at least**

To start a cluster we create a YarnCluster object. We'll create a cluster with 4 workers, each with 4 GB of memory and 2 cores.

```
In [1]: from dask_yarn import YarnCluster  
In [2]: cluster = YarnCluster(environment='environment.tar.gz',  
...:                          worker_vcores=2,  
...:                          worker_memory='4GB',  
...:                          n_workers=4)
```

**Dask+Yarn  
(for example)**

Next we connect to the cluster by creating a `dask.distributed.Client`.

```
In [3]: from dask.distributed import Client  
In [4]: client = Client(cluster)  
In [5]: client  
Out[5]: <Client: scheduler='tcp://172.18.0.2:36217' processes=4 cores=8>
```

# “delayed”

- One more very esoteric use case
  - strips-off most of Dask advantages
  - when the usual data structures cannot be used
  - user has some control on the parallelisation
  - useful for strange data structures and for complex inter-process dependencies
- Small demo using the iqoqo sdk demonstrates the point but
  - we have to force no dependencies
  - We have to force non-Dask data structures
- This approach has no advantage on simply using iqoqomp and it does not exploit any of the main features of Dask.



# “delayed”

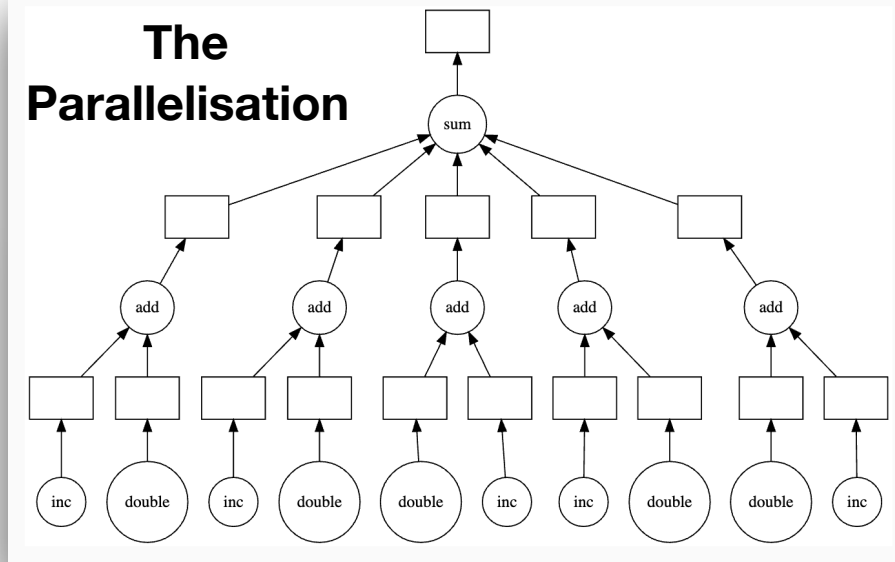
## dask\_func.py (just functions)

```
1 from iqqo import iqqo_job
2
3 @iqqo_job
4 def inc(x):
5     return x + 1
6
7 @iqqo_job
8 def double(x):
9     return x + 2
10
11 @iqqo_job
12 def add(x, y):
13     return x + y
14
```

## dask\_delayed.py (steering)

```
1 import dask
2 from dask import compute, delayed
3 from dask_func import inc, double, add
4 from iqqo import iqqo_job
5
6 data = range(5)
7
8 output = []
9 for x in data:
10     a = dask.delayed(inc)(x)
11     b = dask.delayed(double)(x)
12     c = dask.delayed(add)(a, b)
13     output.append(c)
14
15 total = dask.delayed(sum)(output)
16 print(total.compute())
17
```

## The Parallelisation



**The dependencies won't work in this example on ANY (remote) cluster**

	double1563392456.4703171	2s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago				
	inc1563392456.4712079	3s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago				
	inc1563392456.470722	2s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago				
	inc1563392456.47194	3s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago				
	double1563392456.471457	3s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago				
	double1563392456.472013	3s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago				

Jobs					New job
This page shows all your jobs. Click job name for full details.					
All	Pending	Running	Done	Errors	
add1563392514.219742			2s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago					
			2s	<a href="#">Results</a>	...
add1563392514.142106			2s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago					
			3s	<a href="#">Results</a>	...
add1563392512.683824			3s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago					
			3s	<a href="#">Results</a>	...
add1563392512.710143			3s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago					
			1s	<a href="#">Results</a>	...
add1563392511.581646			1s	<a href="#">Results</a>	...
1/1 Tasks done • created 9 hours ago					